

Churn Case

Mette

2024-03-05

Table of contents

1	Introduction	2
2	Analysis	2
2.1	Business Understanding	2
2.2	Data Understanding	3
2.2.1	Reading libraries	3
2.2.2	Importing data	3
2.3	Data Preparation	5
2.3.1	Cleaning data	5
2.4	Modeling	8
2.4.1	Cost Assessment	8
2.5	Evaluation	16
2.6	Deployment	17
3	Conslusion	17



1

Introduction

This case focuses on customer churn (also known as customer attrition or customer turnover). Customer churn is interesting because it is usually much cheaper to retain existing customers than to acquire new ones. Instead of focusing on each individual customer, we will attempt to build a predictive model that can help us decide which customers we should focus our retention efforts on.

2 Analysis

I follow the CRISP-DM (Cross-Industry Standard Process for Data Mining) framework in my data mining projects, guiding me through six phases: Business Understanding, Data Understanding, Data Preparation, Modeling, Evaluation, and Deployment. This structured approach ensures I effectively extract insights and apply data science.

2.1 Business Understanding

Understanding customer churn is critical for banks as it aids in cost reduction by prioritizing the retention of existing customers over acquiring new ones. This not only helps in stabilizing revenue but also enhances customer satisfaction by addressing their specific needs and concerns. By gaining insights into churn patterns, banks can develop targeted strategies, optimize resource allocation, and gain a competitive edge in the market.

2.2 Data Understanding

2.2.1 Reading libraries

We will start by loading the nessesary libraries

```
pacman::p_load("tidyverse", "magrittr", "nycflights13", "gapminder",
               "Lahman", "maps", "lubridate", "pryr", "hms", "hexbin",
               "feather", "htmlwidgets", "broom", "pander", "modelr",
               "XML", "httr", "jsonlite", "lubridate", "microbenchmark",
               "splines", "ISLR2", "MASS", "testthat", "caret",
               "RSQLite", "class", "babynames", "nasaweather", "pls",
               "fueleconomy", "viridis", "boot", "devtools", "tree", "leaps",
               "glmnet", "gam", "akima", "factoextra", "randomForest", "gbm",
               "ggrepel", "GGally", "fmsb", "sjPlot", "rcompanion", "DT")
```

2.2.2 Importing data

The dataset we are going to work with will be imported and investigated.

```
#bank_churn <- read.csv("Churn_Modelling.csv")
bank_churn <- read.csv("C:/Users/mette/OneDrive/Skrivebord/PB dataanalyse/Programmering og s
```

```
#tjekker data og klasser
str(bank_churn)
```

```
#vi skaber en interaktiv tabel

datatable(bank_churn, caption = htmltools::tags$caption(
  style = 'caption-side: top; text-align: center;',
  'Table 1: ', htmltools::em('Bank data '))
)
```

Show entries

Search:

Table 1: *Bank data*

RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts
1	1	15634602	Hargrave	619	France	Female	42	2	0
2	2	15647311	Hill	608	Spain	Female	41	1	83807.86
3	3	15619304	Onio	502	France	Female	42	8	159660.8
4	4	15701354	Boni	699	France	Female	39	1	0
5	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82
6	6	15574012	Chu	645	Spain	Male	44	8	113755.78
7	7	15592531	Bartlett	822	France	Male	50	7	0
8	8	15656148	Obinna	376	Germany	Female	29	4	115046.74
9	9	15792365	He	501	France	Male	44	4	142051.07
10	10	15592389	H?	684	France	Male	27	2	134603.88

Showing 1 to 10 of 10,000 entries

Previous

2

3

4

5

...

1,000

Next

2.3 Data Preparation

2.3.1 Cleaning data

First we are checking for missing values. There are no missing values in the TotalCharges variable. We are imputing these values with the mean value, since the quantity is low. After this we are changing all the character class variables to factors for later statistical analysis. Finally the churn dataset is again changed to a dataframe and the CustomerID variable is removed.

```
# Beregn antallet af missing værdier i hver kolonne. No missing
bank_churn %>% purrr::map(~ sum(is.na(.)))

summary(bank_churn)

# There are no missing values, so we can proceed.

# the relevant variables are converted into factors.

# Type = factor and integers-----

bank_churn_fact <- bank_churn %>%
  mutate_if(is.character, as.factor) %>%
  mutate_if(is.integer, as.factor)

bank_churn_fact$CreditScore <- as.integer(bank_churn_fact$CreditScore)
bank_churn_fact$Age <- as.integer(bank_churn_fact$Age)
bank_churn_fact$Tenure <- as.integer(bank_churn_fact$Tenure)
str(bank_churn_fact)

# Normalisering -----

# Det er ikke nødvendigt at normalisere data i forbindelse med de statistiske
# modeller, som vi skal køre her. Der er forskellige typer af normalisering.
# Vi ser her på følgende:

normalize <- function(x) {
  ((x-min(x))/(max(x)-min(x)))
}

bank_churn_fact <- bank_churn_fact %>%
  mutate_if(is.numeric, normalize)
```

```

bank_churn_fact <- bank_churn_fact %>%
  dplyr::select(Exited, everything())

glimpse(bank_churn_fact)

numeric_columns <- sapply(bank_churn_fact, is.numeric)

# Konverter numeriske variable fra dbl til int
#bank_churn_fact[numeric_columns] <- lapply(bank_churn_fact[numeric_columns], as.integer)

# Fravalg af customerID, Efternavn og ID -----

bank_churn_fact <- bank_churn_fact %>%
  dplyr::select(-RowNumber, -CustomerId, -Surname)

names(bank_churn_fact)

#Endre variablen Exited til Churn

bank_churn_fact <- bank_churn_fact %>%
  rename(Churn = Exited)

bank_churn_fact$Churn <- ifelse(bank_churn_fact$Churn == 1, "Yes", "No")
bank_churn_fact$Churn <- as.factor(bank_churn_fact$Churn)
str((bank_churn_fact))

```

This creates a dataset with 10.000 observations, that can be investigated and is ready for analysis.

```

# Install DT package
install.packages("DT")

# Load DT package
library(DT)

library(DT)






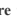

datatable(bank_churn_fact,
  options = list(pageLength = 5, autoWidth = TRUE),
  caption = 'Table 1: Telecommunication data')

```

Show entries

Search:

Table 1: Telecommunication data

	Churn 	CreditScore 	Geography 	Gender 	Age 	Tenure 	Balance 	NumOfPr
1	Yes	0.4967320261437909	France	Female	0.3478260869565217	0.2	0	1
2	No	0.4727668845315904	Spain	Female	0.3333333333333333	0.1	0.334031478677259	1
3	Yes	0.2418300653594771	France	Female	0.3478260869565217	0.8	0.6363571759354565	3
4	No	0.6710239651416122	France	Female	0.3043478260869565	0.1	0	2
5	No	1	Spain	Female	0.3623188405797101	0.2	0.5002462155052675	1

Showing 1 to 5 of 10,000 entries

Previous 2 3 4 5 ... 2,000 Next

2.4 Modeling

Training and test data

```
#træningsdata og testdata

# Vi bruger funktionen set.seed, så vi kan reproducere vores resultater
set.seed(5)
# træningsdel og testdel:
intrain <- createDataPartition(y=bank_churn_fact$Churn,
                               p=0.80, list=FALSE)

# list=FALSE betyder at outputtet bliver en matrix og denne kan bruges
# i koden nedenfor:

train <- bank_churn_fact[intrain,]
test  <- bank_churn_fact[-intrain,]
```

2.4.1 Cost Assessment

The cost assessment is significant when deciding on the threshold in connection with, for example, logistic regression. The relative cost of the different errors that can be made affects where it is optimal to place the threshold. Optimal in terms of reducing costs. When we do not have any a priori knowledge about the relative costs, we use a 50/50 split. But in this example, the situation is different. It does not cost the same to commit the different errors.

#		Vil churne		vil ikke churne
# predikte churne		TP		FP
# predikte ikke churne		FN		TN

Customer acquisition \$200 (<https://www.reviewtrackers.com/blog/customer-acquisition-banking/#:~:text=How%20much%20does%20it%20cost%20a%20bank%20to,%2866%25%29%2C%20and%20SEO%20%2F%20pay-per-click%20%28PPC%29%20advertising%20%2865%25%29.>) - the cost for a false negative (FN) prediction That is, predicting that a customer is satisfied when in reality they churn. #Customer retention \$40 (Source: Bain & Company, “The Value of Online Customer Loyalty in Retail Banking,” 2016.) - the cost of a false positive (FP) That is, predicting that a customer will churn when in reality the customer was satisfied, and a true positive (TP) that is, correctly predicting dissatisfied customers. Correctly predicted true negatives (TN) cost nothing. That is, correctly predicting that a customer is satisfied.

The situation until now is that the company assumes that no customers churn. Thus the model is that nobody churns. In other words, the company predicts a NO for every customer in the test dataset ,i.e., $TP=FP=0$. We only need to calculate FN (since TN does not cost anything) ourselves.

The total savings from the new model based on a customer base of 10.000 customers:

```
FN_omk <- 200

TP_omk <- 40

FP_omk <- TP_omk

TN_omk <- 0
```

Calculation of the current model's costs based on the above information. Since TN is free, we only need to calculate FN on the test dataset. Remember, that all are expected not to churn, but a certain portion will churn.

```
test$no_churn <- "No"
FN_simple <- table(test$Churn, test$no_churn)[2]/(table(test$Churn, test$no_churn)[1]+
                                                    table(test$Churn, test$no_churn)[2])

omkostninger_mavefornehmelse <- FN_omk*FN_simple # per kunde

test <- dplyr::select(test, -no_churn) # We dont need no churn anymore.
```

In order to evaluate the relevant variables we will perform Lasso modelling

```
## results: "hide"
## output: false

bank_churn_lasso <- bank_churn

bank_churn_lasso <- bank_churn_lasso %>%
  dplyr::select(-RowNumber, -CustomerId, -Surname)

names(bank_churn_fact)

#Endre variablen Exited til Churn
```

```

bank_churn_lasso <- bank_churn_lasso %>%
  dplyr::rename(Churn = Exited)

str(bank_churn_lasso)

x <- model.matrix(Churn ~ ., bank_churn_lasso)[, -1]
y <- bank_churn_lasso$Churn

grid <- 10^seq(10, -2, length = 100)
lasso.mod <- glmnet(x, y, alpha = 1, lambda = grid)

coef(lasso.mod)
dim(coef(lasso.mod))

names(bank_churn_lasso)

set.seed(5)
train <- sample(1:nrow(x), nrow(x)*2/3)
test <- (-train)
y.test <- y[test]

set.seed(5)
cv.out <- cv.glmnet(x[train, ], y[train], alpha = 1)
plot(cv.out)

```

```

bestlam <- cv.out$lambda.min
bestlam # optimale
cv.out$lambda
cv.out$lambda.1se

log(bestlam)

lasso.pred <- predict(lasso.mod, s = bestlam, newx = x[test, ])
mse_lasso <- mean((lasso.pred - y.test)^2)

mse_lasso

```

```

out <- glmnet(x, y, alpha = 1)
lasso.coef <- predict(out, type = "coefficients", s = bestlam)[1:8, ]

#print koeficienterne
lasso.coef

```

Based on the lassomodel, the following variables are relevant

```

#| output: true
#| result: true

lasso.coef[lasso.coef != 0]

```

Now we will perform logistic regression

```

#training and test data

train <- bank_churn_fact[intrain,]
test <- bank_churn_fact[-intrain,]

names(test)

outcome <- "Churn"

#Vi så i lasso regression, hvilke variabler der ikke havde relevans, disse ekskluderes
variables <- c( ".", "NumOfProducts", "HasCrCard", "IsActiveMember",
               "EstimatedSalary")

f <- as.formula(paste(outcome,
                      paste(variables, collapse = " - "), sep = " ~ "))

# Vi fitter en logistisk regressionsmodel:

fit_logit <- glm(f, data=train, family = "binomial")

```

```

# Foretage prædiktioner på testsættet og gemmer dem i et objekt, churn_probs:

churn_probs <- predict(fit_logit, test, type = "response")

head(churn_probs)

# Kan vi gøre det bedre end den simple model (og modellen med 50/50 split)
# Loop:

thresh <- seq(0.01, 1.0, length = 100)
omk <- rep(0, length(thresh))

for (i in 1:length(thresh)) {
  glm.pred <- rep("No", length(churn_probs))
  glm.pred[churn_probs>thresh[i]] <- "Yes"
  glm.pred <- as.factor(glm.pred)
  x <- confusionMatrix(glm.pred, test$Churn, positive = "Yes")
  total <- x$table[1] + x$table[2] + x$table[3] + x$table[4]
  TN <- x$table[1]/total
  FP <- x$table[2]/total
  FN <- x$table[3]/total
  TP <- x$table[4]/total
  omk[i] <- FN*FN_omk + TP*TP_omk + FP*FP_omk + TN*0
}

```

We repeat the above procedure, but only for a model where threshold = 0.5. We call this model `simple_model`, and we compare it with different thresholds.

```

glm.pred <- rep("No", length(churn_probs))
glm.pred[churn_probs>0.5] <- "Yes"
glm.pred <- as.factor(glm.pred)
x <- confusionMatrix(glm.pred, test$Churn, positive = "Yes")
total <- x$table[1] + x$table[2] + x$table[3] + x$table[4]
TN <- x$table[1]/total
FP <- x$table[2]/total
FN <- x$table[3]/total
TP <- x$table[4]/total
omk_simple <- FN*FN_omk + TP*TP_omk + FP*FP_omk + TN*0

```

We visualize the results to be able to see in a single way how much the costs depend on the thresholds.

```

model <- c(rep("optimized", 100), "simple")
cost_thresh <- c(omk, omk_simple)
thresh_plot <- c(thresh, 0.5)

```

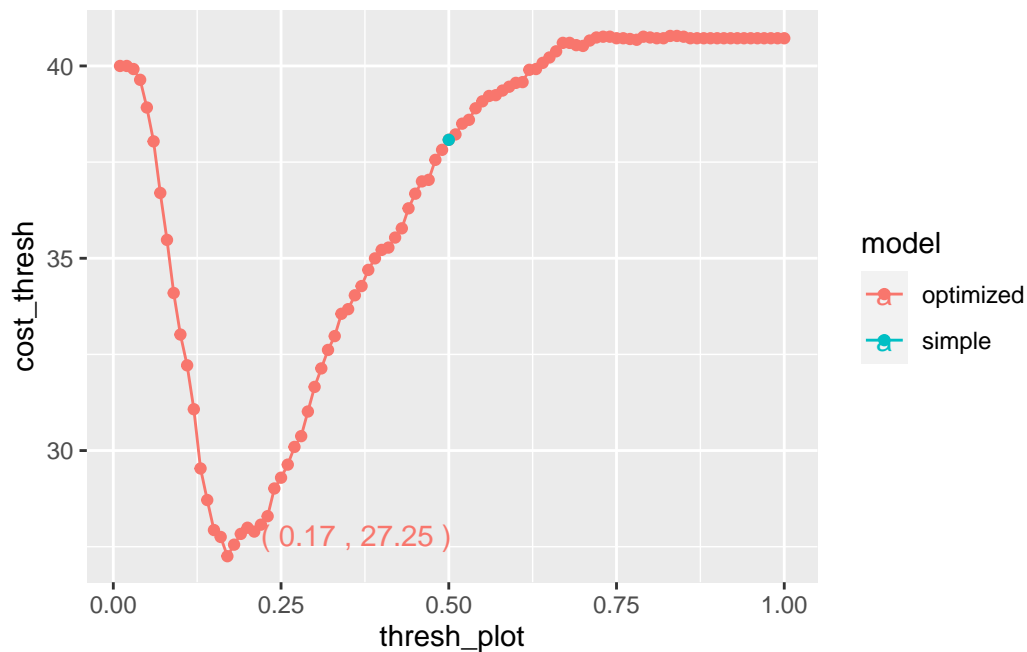
The visualization shows that the optimum threshold is 0,17.

```

dataII <- data.frame(
  model,
  cost_thresh,
  thresh_plot
)

ggplot(dataII, aes(x = thresh_plot, y = cost_thresh, group = model, colour = model)) +
  geom_line() +
  geom_point() +
  geom_text(data = subset(dataII, cost_thresh == min(cost_thresh)),
    aes(label = paste("(", round(thresh_plot, 2), ",", round(cost_thresh, 2), ")"),
      x = thresh_plot + 0.05, y = cost_thresh),
    vjust = -0.5, hjust = 0)

```



Note that we can find an optimum with a different threshold than 0.50.

Calculates the saved costs of the optimized model (threshold=0.17) compared to the baseline model (threshold=0.5).”

```
besparelse_pr_kunde <- omk_simple - min(omk)
besparelse_pr_kunde*10000
```

Calculates the saved costs of the optimized model (threshold=0.17) compared to one’s own calculations.”

```
besparelse_maveforneemmelse <- omkostninger_maveforneemmelse - min(omk)
besparelse_maveforneemmelse*10000
```

We can clearly outperform the existing model that the company uses (which assumes no churn). We can also clearly outperform the common 50/50 split.

Can we do even better with an alternative model: Let’s use linear and quadratic discriminant analysis and quadratic discriminant analysis, and we calculate again the total cost savings, and compare them with the best logistic regression:

2.4.1.1 Linear Discriminant Analysis (LDA)

LDA is a method used for classification and dimensionality reduction. It finds the best linear combination of features to separate different classes in the dataset. LDA works well when classes are distinct and follows normal distributions with equal covariance. It’s useful for understanding which features are most important for classification. However, it may not perform well with overlapping classes or outliers. Overall, LDA is effective for classification tasks with well-separated classes and can provide valuable insights into the data structure.

```
# lda
lda.fit <- lda(f, data=train)

lda.pred <- data.frame(predict(lda.fit, test))
omk_lda <- rep(0,length(thresh))
thresh <- seq(0.01, 1.0, length = 100)

for (i in 1:length(thresh)) {
  glm.pred <- rep("No", length(lda.pred$posterior.Yes)) # laver en vektor med No
  glm.pred[lda.pred$posterior.Yes > thresh[i]] <- "Yes" # ændrer værdien af glm.pred, hvis p
  glm.pred <- as.factor(glm.pred) # sikrer, at variablen er en faktor
  x <- confusionMatrix(glm.pred, test$Churn, positive = "Yes")
  total <- x$table[1] + x$table[2] + x$table[3] + x$table[4]
```

```

TN <- x$table[1]/total
FP <- x$table[2]/total
FN <- x$table[3]/total
TP <- x$table[4]/total
omk_lda[i] <- FN*FN_omk + TP*TP_omk + FP*FP_omk + TN*0
}

```

After running the lda analysis we can compare the savings to the prior models

```
10000*(min(omk_lda) - min(omk))
```

```
[1] -400.2001
```

```
10000*(min(omk_lda) - omk_simple)
```

```
[1] -108654.3
```

```
10000*(min(omk_lda) - omkostninger_maveforneemelse)
```

```
[1] -135067.5
```

2.4.1.2 Quadratic Discriminant Analysis (QDA)

QDA is similar to LDA but allows for different covariance matrices for each class, making it more flexible in capturing complex relationships between features. It works by estimating separate covariance matrices for each class, which can better capture non-linear decision boundaries. QDA is beneficial when classes have different variances or when the decision boundary is non-linear. However, it requires more parameters to estimate compared to LDA and may overfit with small datasets. Overall, QDA is useful for classification tasks with non-linear decision boundaries and varying variances between classes.

```

#qda

qda.fit <- qda(f, data=train)
qda.pred <- data.frame(predict(qda.fit, test))
omk_qda <- rep(0,length(thresh))
thresh <- seq(0.01, 1.0, length = 100)

```

```

for (i in 1:length(thresh)) {
  glm.pred <- rep("No", length(qda.pred$posterior.Yes)) # laver en vektor med No
  glm.pred[qda.pred$posterior.Yes > thresh[i]] <- "Yes" # ændrer værdien af glm-pred, hvis p
  glm.pred <- as.factor(glm.pred) # sikrer at variabelen er en faktor
  x <- confusionMatrix(glm.pred, test$Churn, positive = "Yes")
  total <- x$table[1] + x$table[2] + x$table[3] + x$table[4]
  TN <- x$table[1]/total
  FP <- x$table[2]/total
  FN <- x$table[3]/total
  TP <- x$table[4]/total
  omk_qda[i] <- FN*FN_omk + TP*TP_omk + FP*FP_omk + TN*0
}

```

After running the lda analysis we can compare the savings to the prior models

```
10000*(min(omk_qda) - min(omk))
```

```
[1] 15007.5
```

```
10000*(min(omk_qda) - min(omk_lda))
```

```
[1] 15407.7
```

```
10000*(min(omk_qda) - omk_simple)
```

```
[1] -93246.62
```

```
10000*(min(omk_qda) - omkostninger_maveforneemelse)
```

```
[1] -119659.8
```

2.5 Evaluation

We can now asses the performance of the models. Ss we can se we are able to save money by focusing on datadriven decisions.

2.6 Deployment

The code is scalable since it can be adapted to changes in costs. If we assume that retention cost is double, we only have to change one parameter.

```
TP_omk <- 80
```

And run the code again.

Furthermore once models are trained and evaluated, they're deployed for real-world use. Here's how:

1. **Packaging:** Bundle the trained model with preprocessing steps for easy deployment.
2. **Integration:** Integrate the model into existing systems, ensuring compatibility.
3. **Optimization:** Optimize for scalability and performance to handle real-time requests.
4. **Monitoring:** Monitor model performance over time and update regularly.
5. **Security:** Ensure compliance with security regulations and protect data privacy.
6. **Documentation:** Provide user-friendly documentation and support resources.
7. **Feedback:** Gather feedback for continuous improvement and iteration.

3 Conclusion

Based on the analysis, it is recommended that the bank focuses its retention efforts on customers identified as high-risk churners by the predictive models. This targeted approach can help optimize resource allocation and improve overall customer retention strategies.

In conclusion, the analysis highlights the importance of leveraging data-driven approaches to understand and address customer churn effectively. By implementing the recommended strategies, the bank can enhance customer satisfaction, reduce churn rates, and ultimately, drive long-term business growth and profitability.